

- The process of converting or mapping data from the initial "raw" form into another format, to make it ready for further analysis.
  - It is also known as Data Cleaning and Data Wrangling.
1. Identify, Evaluate and Count missing data
  2. Deal with missing data
  3. Correct the Data Format and Standardize the Data
  4. Normalize the Data (centering/scaling)
  5. Data Binning
  6. Turn Categorical values into Numeric values

**Roll No:- 22354 Name:- Malhar Choure**

**Performance Date: 30/3/2022**

**Submission date: 13/4/2022**

- A -> To **create** cell **above**
- B -> To **create** Cell **below**
- D D -> For **deleting** the cell
- M -> To **markdown** the Cell
- Y -> For **code** the cell
- Z -> To **undo** the deleted cell

You can find the "Automobile Dataset" from the following link:

<https://archive.ics.uci.edu/ml/machine-learning-databases/autos/imports-85.data>.

*# Import the libraries pandas and matplotlib*

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

First, we assign the URL of the dataset to "filename".

Note: This file does not have column headers, which we need to assign.

```
filename = 'https://archive.ics.uci.edu/ml/machine-learning-databases/autos/imports-85.data'
```

Then, we create a Python list headers containing name of headers.

```
headers = ["symboling", "normalized-losses", "make", "fuel-type", "aspiration", "num-of-doors", "body-style", "drive-wheels", "engine-location", "wheel-base", "length", "width", "height", "curb-weight", "engine-type", "num-of-cylinders", "engine-size", "fuel-system", "bore", "stroke", "compression-ratio", "horsepower", "peak-rpm", "city-mpg", "highway-mpg", "price"]
```

Use the Pandas method `read_csv()` to load the data from the web address. Set the parameter "names" equal to the Python list "headers".

```
df = pd.read_csv(filename, names = headers)
```

Use the method `head()` to display the first five rows of the dataframe.

```
# To see what the data set looks like, we'll use the head() method.  
df.head()
```

- Missing values occur when no data value is stored for a variable(feature) in an observation.
- Could be represented as ?, NA, 0 or just a blank cell.

### *Convert "?" to NaN*

In the car dataset, missing data comes with the question mark "?". We replace "?" with NaN (Not a Number), Python's default missing value marker for reasons of computational speed and convenience. Here we use the function: `dataframe.replace(A, B, inplace = True)` to replace A by B.

```
# replace "?" to NaN
```

```
df.replace("?", np.nan, inplace = True) # Question: explain the  
meaning of "inplace = True"  
df.head(5)
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-
0	of-doors \	3	NaN	alfa-romero	gas	std
1	two	3	NaN	alfa-romero	gas	std
2	two	1	NaN	alfa-romero	gas	std
3	four	2	164	audi	gas	std
4	four	2	164	audi	gas	std

	body-style	drive-wheels	engine-location	wheel-base	...	engine-
0	convertible	rwd	front	88.6	...	
1	convertible	rwd	front	88.6	...	
2	hatchback	rwd	front	94.5	...	
3	sedan	fwd	front	99.8	...	
4	sedan	4wd	front	99.4	...	

136

	fuel-system	bore	stroke	compression-ratio	horsepower	peak-rpm
city-mpg \						
0	mpfi	3.47	2.68	9.0	111	5000
21						
1	mpfi	3.47	2.68	9.0	111	5000
21						
2	mpfi	2.68	3.47	9.0	154	5000
19						
3	mpfi	3.19	3.40	10.0	102	5500
24						
4	mpfi	3.19	3.40	8.0	115	5500
18						

	highway-mpg	price
0	27	13495
1	27	16500
2	26	16500
3	30	13950
4	22	17450

[5 rows x 26 columns]

The missing values (NaN) are converted by default. We use the following functions to identify these missing values. There are two methods to detect missing data:

The output is a boolean value indicating whether the value that is passed into the argument is in fact missing data.

```
missing_data = df.isnull()
missing_data.head(5)
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-
doors \						
0	False	True	False	False	False	
False						
1	False	True	False	False	False	
False						
2	False	True	False	False	False	
False						
3	False	False	False	False	False	
False						
4	False	False	False	False	False	
False						

	body-style	drive-wheels	engine-location	wheel-base	...	engine-
size \						
0	False	False	False	False	...	
False						

```

1      False      False      False      False      False ...
False
2      False      False      False      False      False ...
False
3      False      False      False      False      False ...
False
4      False      False      False      False      False ...
False

```

	fuel-system	bore	stroke	compression-ratio	horsepower	peak-rpm
0	False	False	False	False	False	False
1	False	False	False	False	False	False
2	False	False	False	False	False	False
3	False	False	False	False	False	False
4	False	False	False	False	False	False

	city-mpg	highway-mpg	price
0	False	False	False
1	False	False	False
2	False	False	False
3	False	False	False
4	False	False	False

[5 rows x 26 columns]

"True" means the value is a missing value while "False" means the value is not a missing value.

Using a for loop in Python, we can quickly figure out the number of missing values in each column. As mentioned above, "True" represents a missing value and "False" means the value is present in the dataset. In the body of the for loop the method ".value\_counts()" counts the number of "True" values.

```

for column in missing_data.columns.values.tolist():
    print(column)
    print (missing_data[column].value_counts())
    print("")

```

```

symboling
False      205
Name: symboling, dtype: int64

```

```

normalized-losses
False      164

```

True 41  
Name: normalized-losses, dtype: int64

make  
False 205  
Name: make, dtype: int64

fuel-type  
False 205  
Name: fuel-type, dtype: int64

aspiration  
False 205  
Name: aspiration, dtype: int64

num-of-doors  
False 203  
True 2  
Name: num-of-doors, dtype: int64

body-style  
False 205  
Name: body-style, dtype: int64

drive-wheels  
False 205  
Name: drive-wheels, dtype: int64

engine-location  
False 205  
Name: engine-location, dtype: int64

wheel-base  
False 205  
Name: wheel-base, dtype: int64

length  
False 205  
Name: length, dtype: int64

width  
False 205  
Name: width, dtype: int64

height  
False 205  
Name: height, dtype: int64

curb-weight

False 205  
Name: curb-weight, dtype: int64

engine-type  
False 205  
Name: engine-type, dtype: int64

num-of-cylinders  
False 205  
Name: num-of-cylinders, dtype: int64

engine-size  
False 205  
Name: engine-size, dtype: int64

fuel-system  
False 205  
Name: fuel-system, dtype: int64

bore  
False 201  
True 4  
Name: bore, dtype: int64

stroke  
False 201  
True 4  
Name: stroke, dtype: int64

compression-ratio  
False 205  
Name: compression-ratio, dtype: int64

horsepower  
False 203  
True 2  
Name: horsepower, dtype: int64

peak-rpm  
False 203  
True 2  
Name: peak-rpm, dtype: int64

city-mpg  
False 205  
Name: city-mpg, dtype: int64

highway-mpg  
False 205

Name: highway-mpg, dtype: int64

```
price
False    201
True      4
Name: price, dtype: int64
```

Based on the summary above, each column has 205 rows of data and seven of the columns containing missing data:

- **Check with the data collection source**
- **Replace the missing values**
  - replace it with an average (of similar data points)
  - replace it by frequency
  - replace it based on other functions
- **Drop the missing values**
  - drop the variable (column)
  - drop the data entry (row)
- **Leave it as missing data**

Use `dataframe.replace(missing_data, new_data)`

```
avg_norm_loss = df["normalized-losses"].astype("float").mean(axis=0)
print("Average of normalized-losses:", avg_norm_loss)
```

Average of normalized-losses: 122.0

```
df["normalized-losses"].replace(np.nan, avg_norm_loss, inplace=True)
```

```
df.head(5)
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-
0	three	122.0	alfa-romero	gas	std	two
1	three	122.0	alfa-romero	gas	std	two
2	one	122.0	alfa-romero	gas	std	two
3	two	164	audi	gas	std	four
4	two	164	audi	gas	std	four

	body-style	drive-wheels	engine-location	wheel-base	...	engine-
0	convertible	rwd	front	88.6	...	size
1	convertible	rwd	front	88.6	...	130

```

130
2    hatchback          rwd          front          94.5 ...
152
3          sedan          fwd          front          99.8 ...
109
4          sedan          4wd          front          99.4 ...
136

```

```

      fuel-system  bore  stroke  compression-ratio  horsepower  peak-rpm
city-mpg \
0      mpfi  3.47    2.68             9.0           111       5000
21
1      mpfi  3.47    2.68             9.0           111       5000
21
2      mpfi  2.68    3.47             9.0           154       5000
19
3      mpfi  3.19    3.40            10.0           102       5500
24
4      mpfi  3.19    3.40             8.0           115       5500
18

```

```

      highway-mpg  price
0          27  13495
1          27  16500
2          26  16500
3          30  13950
4          22  17450

```

[5 rows x 26 columns]

```

avg_bore=df['bore'].astype('float').mean(axis=0)
print("Average of bore:", avg_bore)

```

Average of bore: 3.3297512437810957

```

df["bore"].replace(np.nan, avg_bore, inplace=True)

```

```

# Write your code below and press Shift+Enter to execute
#Calculate the mean vaule for "stroke" column
avg_stroke = df["stroke"].astype("float").mean(axis = 0)
print("Average of stroke:", avg_stroke)

```

```

# replace NaN by mean value in "stroke" column
df["stroke"].replace(np.nan, avg_stroke, inplace = True)

```

Average of stroke: 3.2554228855721337

```

avg_horsepower = df['horsepower'].astype('float').mean(axis=0)
print("Average horsepower:", avg_horsepower)

```

Average horsepower: 104.25615763546799



```
df['horsepower'].replace(np.nan, avg_horsepower, inplace=True)
```

```
avg_peakrpm=df['peak-rpm'].astype('float').mean(axis=0)
print("Average peak rpm:", avg_peakrpm)
```

Average peak rpm: 5125.369458128079

```
df['peak-rpm'].replace(np.nan, avg_peakrpm, inplace=True)
```

To see which values are present in a particular column, we can use the ".value\_counts()" method:

```
df['num-of-doors'].value_counts()
```

```
four    114
two      89
Name: num-of-doors, dtype: int64
```

We can see that four doors are the most common type. We can also use the ".idxmax()" method to calculate the most common type automatically:

```
df['num-of-doors'].value_counts().idxmax()
```

```
'four'
```

The replacement procedure is very similar to what we have seen previously:

```
#replace the missing 'num-of-doors' values by the most frequent
df["num-of-doors"].replace(np.nan, "four", inplace=True)
```

```
df.head(10)
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-
of-doors \						
0	3	122.0	alfa-romero	gas	std	
two						
1	3	122.0	alfa-romero	gas	std	
two						
2	1	122.0	alfa-romero	gas	std	
two						
3	2	164	audi	gas	std	
four						
4	2	164	audi	gas	std	
four						
5	2	122.0	audi	gas	std	
two						
6	1	158	audi	gas	std	
four						
7	1	122.0	audi	gas	std	
four						
8	1	158	audi	gas	turbo	
four						
9	0	122.0	audi	gas	turbo	

two

	body-style	drive-wheels	engine-location	wheel-base	...	engine-size \
0	convertible	rwd	front	88.6	...	130
1	convertible	rwd	front	88.6	...	130
2	hatchback	rwd	front	94.5	...	152
3	sedan	fwd	front	99.8	...	109
4	sedan	4wd	front	99.4	...	136
5	sedan	fwd	front	99.8	...	136
6	sedan	fwd	front	105.8	...	136
7	wagon	fwd	front	105.8	...	136
8	sedan	fwd	front	105.8	...	131
9	hatchback	4wd	front	99.5	...	131

	fuel-system	bore	stroke	compression-ratio	horsepower	peak-rpm
0	mpfi	3.47	2.68	9.0	111	5000
21						
1	mpfi	3.47	2.68	9.0	111	5000
21						
2	mpfi	2.68	3.47	9.0	154	5000
19						
3	mpfi	3.19	3.40	10.0	102	5500
24						
4	mpfi	3.19	3.40	8.0	115	5500
18						
5	mpfi	3.19	3.40	8.5	110	5500
19						
6	mpfi	3.19	3.40	8.5	110	5500
19						
7	mpfi	3.19	3.40	8.5	110	5500
19						
8	mpfi	3.13	3.40	8.3	140	5500
17						
9	mpfi	3.13	3.40	7.0	160	5500
16						

	highway-mpg	price
0	27	13495

1	27	16500
2	26	16500
3	30	13950
4	22	17450
5	25	15250
6	25	17710
7	25	18920
8	20	23875
9	22	NaN

[10 rows x 26 columns]

- Use `dataframe.dropna()`
  - `axis= 0` to drop the entire row
  - `axis= 1` to drop the entire column
- Whole columns should be dropped only if most entries in the column are empty. In our dataset, none of the columns are empty enough to drop entirely.
- Drop the whole row: "price": 4 missing data, simply delete the whole row Reason: price is what we want to predict in later experiment. Any data entry without price data cannot be used for prediction; therefore any row now without price data is not useful to us

*# simply drop whole row with NaN in "price" column*

`df.dropna(subset=["price"], axis=0, inplace=True)` *# equivalent to: df = df.dropna(subset= ['price'], axis= 0)*

*# reset index, because we dropped two rows*

`df.reset_index(drop=True, inplace=True)`

`print(df)`

	symboling	normalized-losses	make	fuel-type	aspiration	\
0	3	122.0	alfa-romero	gas	std	
1	3	122.0	alfa-romero	gas	std	
2	1	122.0	alfa-romero	gas	std	
3	2	164	audi	gas	std	
4	2	164	audi	gas	std	
...	...	...	...	...	...	...
196	-1	95	volvo	gas	std	
197	-1	95	volvo	gas	turbo	
198	-1	95	volvo	gas	std	
199	-1	95	volvo	diesel	turbo	
200	-1	95	volvo	gas	turbo	

	num-of-doors	body-style	drive-wheels	engine-location	wheel-base
...	\				
0	two	convertible	rwd	front	88.6
...					

1	two	convertible	rwd	front	88.6
...					
2	two	hatchback	rwd	front	94.5
...					
3	four	sedan	fwd	front	99.8
...					
4	four	sedan	4wd	front	99.4
...					
..	...	...	...	...	...
...					
196	four	sedan	rwd	front	109.1
...					
197	four	sedan	rwd	front	109.1
...					
198	four	sedan	rwd	front	109.1
...					
199	four	sedan	rwd	front	109.1
...					
200	four	sedan	rwd	front	109.1
...					

	engine-size	fuel-system	bore	stroke	compression-ratio
horsepower \					
0	130	mpfi	3.47	2.68	9.0
111					
1	130	mpfi	3.47	2.68	9.0
111					
2	152	mpfi	2.68	3.47	9.0
154					
3	109	mpfi	3.19	3.40	10.0
102					
4	136	mpfi	3.19	3.40	8.0
115					
..	...	...	...	...	...
..					
196	141	mpfi	3.78	3.15	9.5
114					
197	141	mpfi	3.78	3.15	8.7
160					
198	173	mpfi	3.58	2.87	8.8
134					
199	145	idi	3.01	3.40	23.0
106					
200	141	mpfi	3.78	3.15	9.5
114					

	peak-rpm	city-mpg	highway-mpg	price
0	5000	21	27	13495
1	5000	21	27	16500
2	5000	19	26	16500

3	5500	24	30	13950
4	5500	18	22	17450
...	...	...	...	...
196	5400	23	28	16845
197	5300	19	25	19045
198	5500	18	23	21485
199	4800	26	27	22470
200	5400	19	25	22625

[201 rows x 26 columns]

Good! Now, we have a dataset with no missing values.

In this section, we will look at the problem of data with different formats, units and conventions and the pandas methods that help us deal with these issues.

- Data are generally collected from different places and stored in different formats.
- Data formatting and standardization: Bringing (transforming) data into a common standard of expression allow users to make meaningful comparison.
- As a part of data cleaning, formatting ensures the data is consistent and easily understandable.

Steps for Data formatting and standardization

- Correcting the incorrect data types (Data Formatting)
- Applying calculation to an entire column (Data Standardization)

In Pandas, we use:

`df.dtypes`

symboling	int64
normalized-losses	object
make	object
fuel-type	object
aspiration	object
num-of-doors	object
body-style	object
drive-wheels	object
engine-location	object
wheel-base	float64
length	float64
width	float64
height	float64
curb-weight	int64
engine-type	object
num-of-cylinders	object
engine-size	int64
fuel-system	object
bore	object
stroke	object

```

compression-ratio    float64
horsepower            object
peak-rpm              object
city-mpg              int64
highway-mpg           int64
price                 object
dtype: object

```

```
df
```

```

      symboling normalized-losses      make fuel-type aspiration \
0           3          122.0  alfa-romero    gas          std
1           3          122.0  alfa-romero    gas          std
2           1          122.0  alfa-romero    gas          std
3           2           164      audi      gas          std
4           2           164      audi      gas          std
..      ...      ...      ...      ...      ...
196        -1           95    volvo      gas          std
197        -1           95    volvo      gas        turbo
198        -1           95    volvo      gas          std
199        -1           95    volvo  diesel        turbo
200        -1           95    volvo      gas        turbo

```

```

      num-of-doors  body-style drive-wheels engine-location  wheel-base
... \
0           two  convertible          rwd          front      88.6
...
1           two  convertible          rwd          front      88.6
...
2           two   hatchback          rwd          front      94.5
...
3           four      sedan          fwd          front      99.8
...
4           four      sedan          4wd          front      99.4
...
..      ...      ...      ...      ...      ...
...
196        four      sedan          rwd          front     109.1
...
197        four      sedan          rwd          front     109.1
...
198        four      sedan          rwd          front     109.1
...
199        four      sedan          rwd          front     109.1
...
200        four      sedan          rwd          front     109.1
...

```

```

      engine-size  fuel-system  bore  stroke compression-ratio
horsepower \

```

0	130	mpfi	3.47	2.68	9.0
111					
1	130	mpfi	3.47	2.68	9.0
111					
2	152	mpfi	2.68	3.47	9.0
154					
3	109	mpfi	3.19	3.40	10.0
102					
4	136	mpfi	3.19	3.40	8.0
115					
..	...	...	...	...	..
196	141	mpfi	3.78	3.15	9.5
114					
197	141	mpfi	3.78	3.15	8.7
160					
198	173	mpfi	3.58	2.87	8.8
134					
199	145	idi	3.01	3.40	23.0
106					
200	141	mpfi	3.78	3.15	9.5
114					

	peak-rpm	city-mpg	highway-mpg	price
0	5000	21	27	13495
1	5000	21	27	16500
2	5000	19	26	16500
3	5500	24	30	13950
4	5500	18	22	17450
..	...	...	...	...
196	5400	23	28	16845
197	5300	19	25	19045
198	5500	18	23	21485
199	4800	26	27	22470
200	5400	19	25	22625

[201 rows x 26 columns]

```
df[["bore", "stroke"]] = df[["bore", "stroke"]].astype("float")
df[["normalized-losses"]] = df[["normalized-losses"]].astype("int")
df[["price"]] = df[["price"]].astype("float")
df[["peak-rpm"]] = df[["peak-rpm"]].astype("float")
```

df.dtypes

symboling	int64
normalized-losses	int32
make	object
fuel-type	object
aspiration	object
num-of-doors	object

```

body-style          object
drive-wheels        object
engine-location      object
wheel-base          float64
length              float64
width               float64
height              float64
curb-weight          int64
engine-type          object
num-of-cylinders     object
engine-size          int64
fuel-system          object
bore                float64
stroke              float64
compression-ratio    float64
horsepower           object
peak-rpm             float64
city-mpg             int64
highway-mpg          int64
price                float64
dtype: object

```

Wonderful!

Now we have finally obtained the cleaned dataset with no missing values with all data in its proper format.

Example

The formula for unit conversion is:  $L/100km = 235 / mpg$  We can do many mathematical operations directly in Pandas.

```
df.head()
```

```

   symboling  normalized-losses      make fuel-type aspiration \
0           3                122  alfa-romero      gas      std
1           3                122  alfa-romero      gas      std
2           1                122  alfa-romero      gas      std
3           2                164      audi      gas      std
4           2                164      audi      gas      std

```

```

   num-of-doors  body-style drive-wheels engine-location  wheel-
base ... \
0           two  convertible          rwd          front
88.6 ...
1           two  convertible          rwd          front
88.6 ...
2           two   hatchback          rwd          front
94.5 ...
3          four     sedan          fwd          front
99.8 ...

```



4	four	sedan	4wd	front		
99.4	...					
	engine-size	fuel-system	bore	stroke	compression-ratio	horsepower
\						
0	130	mpfi	3.47	2.68	9.0	111
1	130	mpfi	3.47	2.68	9.0	111
2	152	mpfi	2.68	3.47	9.0	154
3	109	mpfi	3.19	3.40	10.0	102
4	136	mpfi	3.19	3.40	8.0	115

	peak-rpm	city-mpg	highway-mpg	price
0	5000.0	21	27	13495.0
1	5000.0	21	27	16500.0
2	5000.0	19	26	16500.0
3	5500.0	24	30	13950.0
4	5500.0	18	22	17450.0

[5 rows x 26 columns]

*# Convert mpg to L/100km by mathematical operation (235 divided by mpg)*

*df['city-L/100km'] = 235/df["city-mpg"] # This will create a new column "city-L/100km"*

*# check your transformed data*

df.head()

	symboling	normalized-losses	make	fuel-type	aspiration	\
0	3	122	alfa-romero	gas	std	
1	3	122	alfa-romero	gas	std	
2	1	122	alfa-romero	gas	std	
3	2	164	audi	gas	std	
4	2	164	audi	gas	std	

	num-of-doors	body-style	drive-wheels	engine-location	wheel-
base	...	\			
0	two	convertible	rwd	front	
88.6	...				
1	two	convertible	rwd	front	
88.6	...				
2	two	hatchback	rwd	front	
94.5	...				
3	four	sedan	fwd	front	
99.8	...				

```

4          four          sedan          4wd          front
99.4  ...

      fuel-system  bore  stroke  compression-ratio  horsepower  peak-rpm
city-mpg \
0          mpfi  3.47   2.68              9.0          111   5000.0
21
1          mpfi  3.47   2.68              9.0          111   5000.0
21
2          mpfi  2.68   3.47              9.0          154   5000.0
19
3          mpfi  3.19   3.40             10.0          102   5500.0
24
4          mpfi  3.19   3.40              8.0          115   5500.0
18

```

```

      highway-mpg      price  city-L/100km
0          27   13495.0      11.190476
1          27   16500.0      11.190476
2          26   16500.0      12.368421
3          30   13950.0       9.791667
4          22   17450.0     13.055556

```

[5 rows x 27 columns]

```

# Write your code below and press Shift+Enter to execute
df["highway-L/100km"]=235/df["highway-mpg"]
df.head(10)

```

```

      symboling  normalized-losses      make fuel-type aspiration \
0          3          122  alfa-romero    gas      std
1          3          122  alfa-romero    gas      std
2          1          122  alfa-romero    gas      std
3          2          164      audi    gas      std
4          2          164      audi    gas      std
5          2          122      audi    gas      std
6          1          158      audi    gas      std
7          1          122      audi    gas      std
8          1          158      audi    gas    turbo
9          2          192      bmw    gas      std

```

```

      num-of-doors  body-style  drive-wheels  engine-location  wheel-
base ... \
0          two  convertible      rwd      front
88.6 ...
1          two  convertible      rwd      front
88.6 ...
2          two   hatchback      rwd      front
94.5 ...
3          four      sedan      fwd      front

```

99.8	...				
4	four	sedan	4wd	front	
99.4	...				
5	two	sedan	fwd	front	
99.8	...				
6	four	sedan	fwd	front	
105.8	...				
7	four	wagon	fwd	front	
105.8	...				
8	four	sedan	fwd	front	
105.8	...				
9	two	sedan	rwd	front	
101.2	...				

	bore	stroke	compression-ratio	horsepower	peak-rpm	city-mpg
highway-mpg \						
0	3.47	2.68	9.0	111	5000.0	21
27						
1	3.47	2.68	9.0	111	5000.0	21
27						
2	2.68	3.47	9.0	154	5000.0	19
26						
3	3.19	3.40	10.0	102	5500.0	24
30						
4	3.19	3.40	8.0	115	5500.0	18
22						
5	3.19	3.40	8.5	110	5500.0	19
25						
6	3.19	3.40	8.5	110	5500.0	19
25						
7	3.19	3.40	8.5	110	5500.0	19
25						
8	3.13	3.40	8.3	140	5500.0	17
20						
9	3.50	2.80	8.8	101	5800.0	23
29						

	price	city-L/100km	highway-L/100km
0	13495.0	11.190476	8.703704
1	16500.0	11.190476	8.703704
2	16500.0	12.368421	9.038462
3	13950.0	9.791667	7.833333
4	17450.0	13.055556	10.681818
5	15250.0	12.368421	9.400000
6	17710.0	12.368421	9.400000
7	18920.0	12.368421	9.400000
8	23875.0	13.823529	11.750000
9	16430.0	10.217391	8.103448

[10 rows x 28 columns]

## Example

Few Methods of normalizing data

1. **Simple feature scaling:**  $x_{new} = \frac{x_{old}}{x_{max}}$
2. **Min-Max:**  $x_{new} = \frac{x_{old} - x_{min}}{x_{max} - x_{min}}$
3. **Z-score:**  $x_{new} = \frac{x_{old} - \mu}{\sigma}$  where  $\mu$  is the mean and  $\sigma$  is the standard deviation of the feature.

```
df['length']  
0      168.8  
1      168.8  
2      171.2  
3      176.6  
4      176.6  
...  
196    188.8  
197    188.8  
198    188.8  
199    188.8  
200    188.8  
Name: length, Length: 201, dtype: float64
```

### 5.1 Simple feature scaling

```
# replace (original value) by (original value)/(maximum value)  
df['length'] = df['length']/df['length'].max()  
df['width'] = df['width']/df['width'].max()  
  
# Write your code below and press Shift+Enter to execute  
df['height'] = df['height']/df['height'].max()  
print(df['height'].max())  
print(df['height'])  
  
1.0  
0      0.816054  
1      0.816054  
2      0.876254  
3      0.908027  
4      0.908027  
...  
196    0.928094  
197    0.928094  
198    0.928094  
199    0.928094
```

```
200      0.928094
Name: height, Length: 201, dtype: float64
```

Here we can see we've normalized "length", "width" and "height" in the range of [0,1].

- **Binning:** Grouping of **values** into **bins** for grouped analysis.
  - Example: we can bin "age" into [0, 5], [6, 10], [11, 15] and so on.
- Converts **numeric** into **categorical** variables.
- Group a **set of numerical values** into a **set of bins**.

Example:

Convert data to correct format:

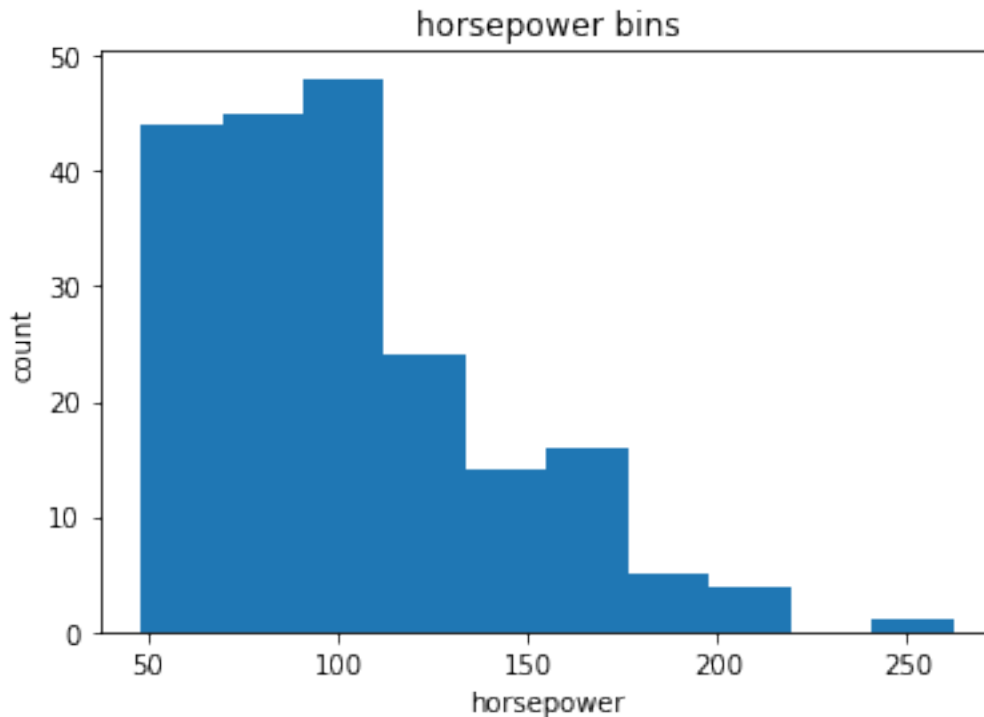
```
df["horsepower"]=df["horsepower"].astype(int, copy=True)
```

Let's plot the histogram of horsepower to see what the distribution of horsepower looks like.

```
%matplotlib inline
import matplotlib as plt
from matplotlib import pyplot
plt.pyplot.hist(df["horsepower"])

# set x/y labels and plot title
plt.pyplot.xlabel("horsepower")
plt.pyplot.ylabel("count")
plt.pyplot.title("horsepower bins")

Text(0.5, 1.0, 'horsepower bins')
```



We build a bin array with a minimum value to a maximum value by using the bandwidth calculated above. The values will determine when one bin ends and another begins.

```
bins = np.linspace(min(df["horsepower"]), max(df["horsepower"]), 4)
bins
array([ 48.          , 119.33333333, 190.66666667, 262.          ])
```

We set group names:

```
group_names = ['Low', 'Medium', 'High']
```

We apply the function "cut" to determine what each value of df['horsepower'] belongs to.

```
df['horsepower-binned'] = pd.cut(df['horsepower'], bins,
labels=group_names, include_lowest=True )
df[['horsepower', 'horsepower-binned']].head(20)
```

	horsepower	horsepower-binned
0	111	Low
1	111	Low
2	154	Medium
3	102	Low
4	115	Low
5	110	Low
6	110	Low
7	110	Low
8	140	Medium

9	101	Low
10	101	Low
11	121	Medium
12	121	Medium
13	121	Medium
14	182	Medium
15	182	Medium
16	182	Medium
17	48	Low
18	70	Low
19	70	Low

Let's see the number of vehicles in each bin:

```
df["horsepower-binned"].value_counts()
```

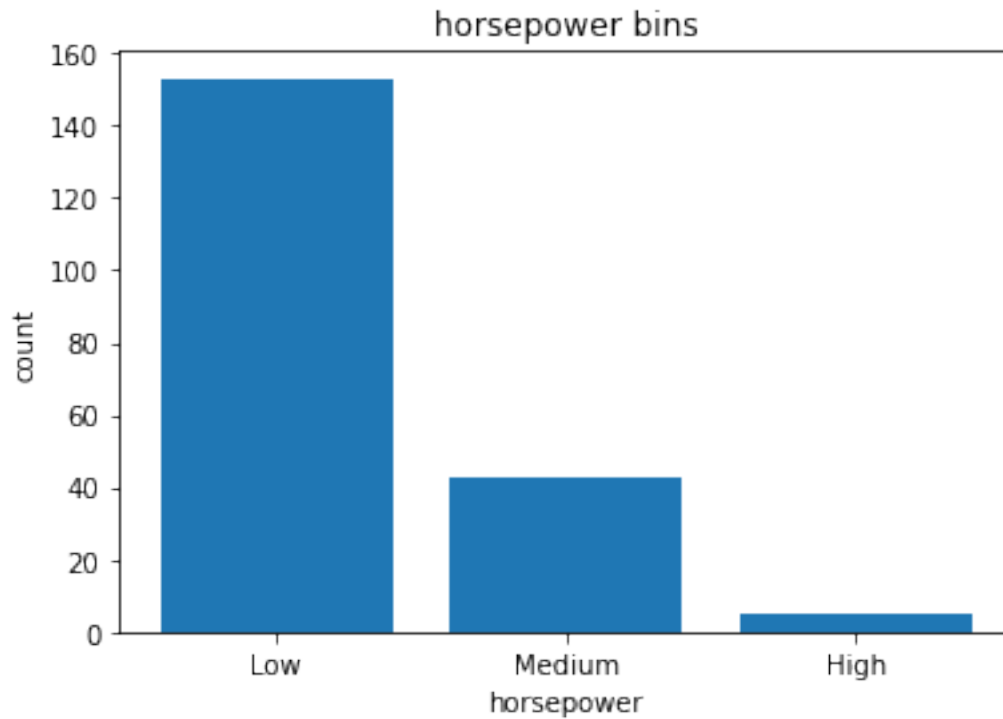
```
Low      153
Medium   43
High      5
Name: horsepower-binned, dtype: int64
```

Let's plot the distribution of each bin:

```
%matplotlib inline
import matplotlib as plt
from matplotlib import pyplot
pyplot.bar(group_names, df["horsepower-binned"].value_counts())

# set x/y labels and plot title
plt.pyplot.xlabel("horsepower")
plt.pyplot.ylabel("count")
plt.pyplot.title("horsepower bins")

Text(0.5, 1.0, 'horsepower bins')
```



Normally, a histogram is used to visualize the distribution of bins we created above.


```
%matplotlib inline
import matplotlib as plt
from matplotlib import pyplot
```

```
# draw histogram of attribute "horsepower" with bins = 3
plt.pyplot.hist(df["horsepower"], bins = 3)
```

```
# set x/y labels and plot title
plt.pyplot.xlabel("horsepower")
plt.pyplot.ylabel("count")
plt.pyplot.title("horsepower bins")
```

```
Text(0.5, 1.0, 'horsepower bins')
```





The plot above shows the binning result for the attribute "horsepower".

What is an indicator variable? An indicator variable (or dummy variable) is a numerical variable used to label categories. They are called 'dummies' because the numbers themselves don't have inherent meaning.

Why we use indicator variables?

Example We see the column "fuel-type" has two unique values: "gas" or "diesel". Regression doesn't understand words, only numbers. To use this attribute in regression analysis, we convert "fuel-type" to indicator variables.

```
df.columns
```

```
Index(['symboling', 'normalized-losses', 'make', 'fuel-type',  
      'aspiration',  
      'num-of-doors', 'body-style', 'drive-wheels', 'engine-  
location',  
      'wheel-base', 'length', 'width', 'height', 'curb-weight',  
      'engine-type',  
      'num-of-cylinders', 'engine-size', 'fuel-system', 'bore',  
      'stroke',  
      'compression-ratio', 'horsepower', 'peak-rpm', 'city-mpg',  
      'highway-mpg', 'price', 'city-L/100km', 'highway-L/100km',  
      'horsepower-binned'],  
      dtype='object')
```

Get the indicator variables and assign it to data frame "dummy\_variable\_1":

```
dummy_variable_1 = pd.get_dummies(df["fuel-type"])
dummy_variable_1.head()
```

```

   diesel  gas
0       0    1
1       0    1
2       0    1
3       0    1
4       0    1

```

Change the column names for clarity:

```
dummy_variable_1.rename(columns={'gas': 'fuel-type-gas',
                                'diesel': 'fuel-type-diesel'}, inplace=True)
dummy_variable_1.head()
```

```

   fuel-type-diesel  fuel-type-gas
0                0                1
1                0                1
2                0                1
3                0                1
4                0                1

```

In the dataframe, column 'fuel-type' has values for 'gas' and 'diesel' as 0s and 1s now.

```
# merge data frame "df" and "dummy_variable_1"
df = pd.concat([df, dummy_variable_1], axis=1)

# drop original column "fuel-type" from "df"
df.drop("fuel-type", axis = 1, inplace=True)

df.head()
```

```

   symboling  normalized-losses  make aspiration num-of-
doors \
0         3           122  alfa-romero      std      two
1         3           122  alfa-romero      std      two
2         1           122  alfa-romero      std      two
3         2           164      audi      std      four
4         2           164      audi      std      four

```

```

   body-style drive-wheels engine-location  wheel-base  length  ...
\
0  convertible          rwd          front      88.6  0.811148  ...
1  convertible          rwd          front      88.6  0.811148  ...

```

2	hatchback	rwd	front	94.5	0.822681	...
3	sedan	fwd	front	99.8	0.848630	...
4	sedan	4wd	front	99.4	0.848630	...

	horsepower	peak-rpm	city-mpg	highway-mpg	price	
city-L/100km	\					
0	111	5000.0	21	27	13495.0	11.190476
1	111	5000.0	21	27	16500.0	11.190476
2	154	5000.0	19	26	16500.0	12.368421
3	102	5500.0	24	30	13950.0	9.791667
4	115	5500.0	18	22	17450.0	13.055556

	highway-L/100km	horsepower-binned	fuel-type-diesel	fuel-type-gas
0	8.703704	Low	0	1
1	8.703704	Low	0	1
2	9.038462	Medium	0	1
3	7.833333	Low	0	1
4	10.681818	Low	0	1

[5 rows x 30 columns]

The last two columns are now the indicator variable representation of the fuel-type variable. They're all 0s and 1s now.

*# Write your code below and press Shift+Enter to execute*  
 print(df["aspiration"])

```

0      std
1      std
2      std
3      std
4      std
...
196    std

```

```
197    turbo
198      std
199    turbo
200    turbo
Name: aspiration, Length: 201, dtype: object
```

Question #5:

Merge the new dataframe to the original dataframe, then drop the column 'aspiration'.

```
# Write your code below and press Shift+Enter to execute
dummy_variable_2= pd.get_dummies(df['aspiration'])
dummy_variable_2.head()
dummy_variable_2.rename(columns={'std':'aspiration_std',
' turbo':'aspiration_turbo'}, inplace=True)
dummy_variable_2.head()
df=pd.concat([df,dummy_variable_2],axis=1)
df.drop("aspiration",axis=1,inplace=True)
```

Save the new csv:

```
df.to_csv('clean_df.csv')
```